



AWS Lambda — Complete Notes

Serverless functions · Run code without managing servers

AWS

Lambda

Model

Serverless

Billing

Per ms

Scaling

Automatic

A reader-friendly, all-in-one reference for AWS Lambda.

Table of Contents

1. What is Lambda
2. Core Concepts
3. The Execution Model
4. Triggers & Event Sources
5. Configuration (Memory, Timeout, etc.)
6. Cold Starts & Concurrency
7. Deployment Packages & Layers
8. Permissions (Execution Role vs Resource Policy)
9. Networking (VPC Access)
10. Versions & Aliases
11. Error Handling & Retries
12. Monitoring & Logging
13. Pricing
14. Common CLI Commands
15. Limits & Best Practices
16. Quick Mental Model
17. Common Interview Questions

1. ⚡ What is Lambda

AWS Lambda runs your code **in response to events** without you provisioning or managing servers. You upload a **function**, Lambda runs it on demand, scales automatically, and you pay only for the **compute time consumed** (per millisecond). No idle servers, no patching the OS.

💡 **Mental model:** Lambda is **Function as a Service (FaaS)**. You own the function code; AWS owns everything below it — servers, scaling, availability, OS.

Common uses: API backends (with API Gateway), event processing (S3/DynamoDB/Kafka), stream processing, scheduled jobs (cron), glue between services, real-time file/data transforms, automation.

2. Core Concepts

Concept	Description
Function	Your code + config (runtime, memory, role, handler).
Handler	The entry-point method Lambda invokes (<code>file.method</code>).
Event	The JSON input passed to the handler.
Context	Runtime info (request id, remaining time, log group).
Runtime	Language environment: Node.js, Python, Java, Go, .NET, Ruby, or custom.
Trigger / Event Source	What invokes the function (S3, API GW, EventBridge...).
Execution Role	IAM role giving the function permissions to call AWS.
Concurrency	Number of simultaneous executions.

3. The Execution Model

1. An **event** arrives from a trigger.
2. Lambda provisions an **execution environment** (micro-VM / Firecracker) if none is warm.
3. It runs **init** code (outside the handler) once per environment.
4. It invokes the **handler** with the event + context.
5. The environment is **frozen** and **reused** for later invocations (warm), then eventually torn down.

Invocation types:

Type	How	Example
Synchronous	Caller waits for the response	API Gateway, ALB, direct invoke
Asynchronous	Event queued; Lambda retries on failure	S3, SNS, EventBridge
Stream / poll-based	Lambda polls and batches records	Kinesis, DynamoDB Streams, SQS, Kafka

 Put expensive setup (DB connections, SDK clients) in the **init** phase (module scope) so warm invocations reuse it.

4. 🛎 Triggers & Event Sources

Lambda integrates with 200+ services. Common triggers:

- **API Gateway / Function URL / ALB** — HTTP APIs and webhooks.
- **S3** — object created/deleted (process uploads).
- **EventBridge / CloudWatch Events** — schedules (cron) and event rules.
- **SNS / SQS** — pub-sub and queue processing.
- **DynamoDB Streams / Kinesis** — stream processing.
- **MSK / self-managed Kafka** — message consumption.
- **Cognito, Step Functions, IoT, CloudFront (Lambda@Edge)** — and many more.

5. ⚙ Configuration (Memory, Timeout, etc.)

Setting	Range / note
Memory	128 MB – 10,240 MB. CPU scales with memory (more memory = more vCPU).
Timeout	1 s – 900 s (15 min) max.
Ephemeral storage (<code>/tmp</code>)	512 MB – 10,240 MB.
Environment variables	Up to 4 KB total; can be KMS-encrypted.
Architecture	x86_64 or arm64 (Graviton) — arm is cheaper/faster for many workloads.

💡 Tuning **memory** is the main performance lever — more memory also raises CPU, often making functions both faster *and* cheaper. Use **Lambda Power Tuning** to find the sweet spot.

6. ❄ Cold Starts & Concurrency

- **Cold start** = the latency to set up a new execution environment (download code, start runtime, run init) before the first invocation. Warm reuse avoids it.
- Bigger packages, VPC attachment, and heavy init increase cold starts. Smaller deps + lazy loading reduce them.
- **Concurrency** = number of simultaneous executions.
 - **Reserved concurrency** — caps (and guarantees) a function's max concurrency.
 - **Provisioned concurrency** — pre-warms N environments → **no cold starts** (extra cost).
- **Account concurrency limit** defaults to **1,000** per region (adjustable). Burst scaling adds environments quickly, then linearly.

7. 📦 Deployment Packages & Layers

- **Package types**: a **.zip archive** (code + deps) or a **container image** (up to 10 GB, OCI).

- **Zip limits:** 50 MB zipped (direct upload), 250 MB unzipped; larger via S3.
- **Layers** — share common libraries/dependencies across functions (up to 5 layers per function). Keeps deployment packages small.
- **Runtimes:** managed (Node, Python, Java, Go, .NET, Ruby) or **custom runtime** via the Runtime API / `provided.al2`.

8. Permissions (Execution Role vs Resource Policy)

Two distinct permission directions — don't confuse them:

	Execution Role	Resource-based Policy
Question	What can the function call ?	Who can invoke the function ?
Type	IAM role assumed by Lambda	Policy attached to the function
Example	Read S3, write DynamoDB, put logs	Allow S3/EventBridge/API GW to invoke

The function always needs `logs:CreateLogGroup/Stream/PutLogEvents` (usually via the `AWSLambdaBasicExecutionRole` managed policy) to write to CloudWatch.

9. Networking (VPC Access)

- By default a Lambda runs in an AWS-managed network with internet access.
- Attach it to a **VPC** (subnets + security group) to reach private resources (RDS, ElastiCache, internal services).
- A VPC-attached Lambda has **no internet by default** — it needs a **NAT Gateway** (in a private subnet) for outbound internet, or **VPC Endpoints** for AWS services.
- Modern Lambda uses **Hyperplane ENIs** (shared), so VPC cold-start penalty is now minimal.

10. Versions & Aliases

- **\$LATEST** — the mutable working version.
- **Published version** — an immutable snapshot (`:1` , `:2` ...).
- **Alias** — a named, movable pointer to a version (`prod` , `staging`).
- Aliases enable **weighted/canary deploys** (e.g. 10% traffic to v2, 90% to v1) and clean rollbacks.

11. Error Handling & Retries

- **Synchronous** invokes: the error returns to the caller; **no automatic retry** (caller decides).
- **Asynchronous** invokes: Lambda **retries twice** (3 total attempts) over time; failures go to a **Dead-Letter Queue (SQS/SNS)** or an **on-failure destination**.

- **Stream/poll** sources (Kinesis/DynamoDB): retry by **batch** until success or record expiry; can block the shard — use bisect-on-error, max retry, and a failure destination.
- **Destinations** (async) route success/failure events to SQS, SNS, Lambda, or EventBridge.

12. 📊 Monitoring & Logging

- **CloudWatch Logs** — automatic; `console.log` / `print` lands in a log group per function.
- **CloudWatch Metrics** — Invocations, Errors, Duration, Throttles, ConcurrentExecutions.
- **AWS X-Ray** — distributed tracing of the function and its downstream calls.
- **Lambda Insights** — enhanced per-invocation system metrics.
- **CloudTrail** — API/management audit.

13. 💰 Pricing

You pay for:

1. **Requests** — per invocation (first 1M/month free, then per million).
2. **Duration** — GB-seconds = (memory ÷ 1024) × execution time, billed per **1 ms**.
3. **Extras** — Provisioned Concurrency (per GB-s reserved), ephemeral storage above 512 MB, data transfer.

🔧 **Cost tips:** right-size **memory** (it sets CPU too) · use **arm64/Graviton** · trim dependencies and init time · avoid idle Provisioned Concurrency · batch records from streams/queues · keep functions short and single-purpose.

14. 📋 Common CLI Commands

```
aws lambda create-function --function-name myfn --runtime python3.12 \
  --handler app.handler --role arn:aws:iam::123:role/lambda-role \
  --zip-file fileb://function.zip

aws lambda invoke --function-name myfn --payload '{"key":"val"}' out.json
aws lambda update-function-code --function-name myfn --zip-file fileb://function.zip
aws lambda update-function-configuration --function-name myfn --memory-size 512 --timeout 30
aws lambda publish-version --function-name myfn
aws lambda create-alias --function-name myfn --name prod --function-version 1
aws lambda list-functions
aws lambda get-function --function-name myfn
```

15. 📌 Limits & Best Practices

Item	Value
Max timeout	15 min (900 s)
Memory	128 MB – 10,240 MB
Ephemeral <code>/tmp</code>	512 MB – 10,240 MB
Default concurrency	1,000 / region (adjustable)
Deployment zip	50 MB (zipped) / 250 MB (unzipped)
Container image	up to 10 GB
Invocation payload	6 MB (sync), 256 KB (async)
Layers per function	5

Best practices: keep functions **small & single-purpose** · do heavy setup in **init** and reuse connections · don't store state in the environment (it's ephemeral) · use **environment variables** (encrypted) for config · least-privilege **execution role** · set sensible **timeouts + DLQ** · use **aliases** for safe deploys · prefer **Provisioned Concurrency** only where cold starts hurt.

16. 🧠 Quick Mental Model

- **Lambda = run code on events, pay per millisecond, zero servers to manage.**
- A **trigger** sends an **event** → Lambda runs your **handler** with an **execution role's** permissions.
- Environments are **reused** (warm) — put setup in **init**; first call may be a **cold start**.
- Tune **memory** (it also sets CPU); cap blast radius with **reserved**, kill cold starts with **provisioned** concurrency.
- **Sync** = caller waits, no retry; **async** = queued + retried + DLQ; **streams** = batched polling.
- Attach to a **VPC** for private resources (needs NAT/endpoints for outbound).
- Ship safely with **versions + aliases**; watch it with **CloudWatch + X-Ray**.

17. ? Common Interview Questions

Rapid-fire questions interviewers ask about Lambda:

- **Q: What is a cold start?** — The latency to set up a new execution environment (download code, start runtime, run init) before the first invocation. Provisioned Concurrency eliminates it.
- **Q: Sync vs async vs stream invocation?** — Sync = caller waits, no auto-retry (API GW). Async = queued, retried twice, DLQ on failure (S3/SNS). Stream/poll = batched records (Kinesis/DynamoDB/SQS).
- **Q: Execution role vs resource-based policy?** — Execution role = what the function **can call**. Resource policy = **who can invoke** the function.
- **Q: How do you tune Lambda performance?** — Increase **memory** (CPU scales with it), trim dependencies, reuse connections via the init phase, use arm64/Graviton.
- **Q: Max timeout and memory?** — 15 minutes; 128 MB–10,240 MB.

- **Q: How to reach a private RDS from Lambda?** — Attach the function to the VPC (subnets + SG); add NAT/endpoints if it also needs outbound internet.
 - **Q: How do versions and aliases help?** — Immutable versions + movable aliases enable canary/weighted deploys and clean rollbacks.
 - **Q: When NOT to use Lambda?** — Long-running (>15 min) jobs, very high steady throughput where always-on compute is cheaper, or workloads needing heavy local state.
-

 Notes compiled as a quick reference for AWS Lambda.